

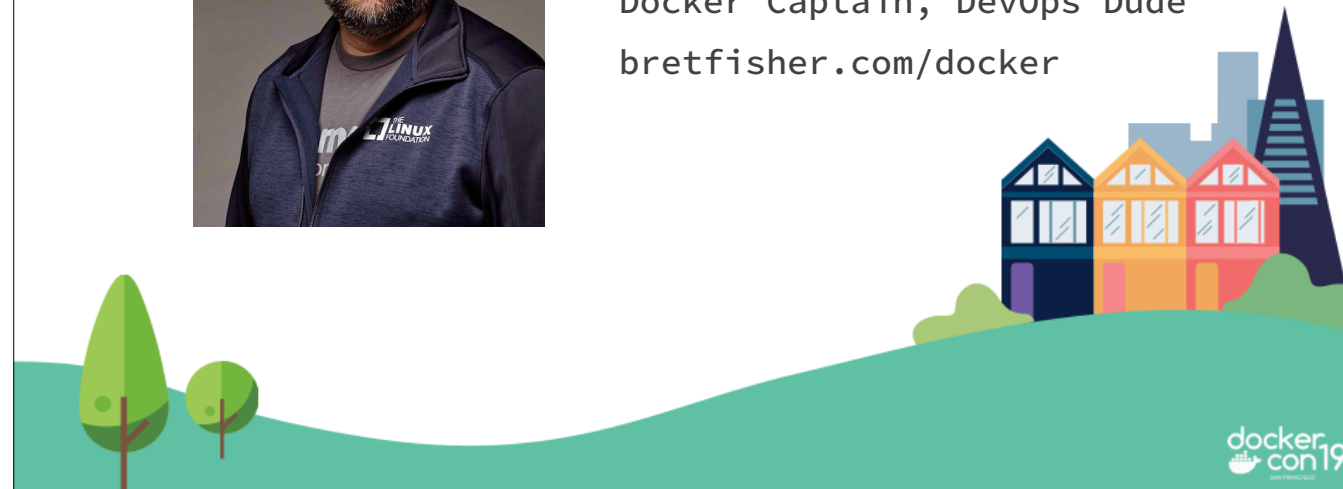


Node.js Rocks in Docker

Bret Fisher

Docker Captain, DevOps Dude

bretfisher.com/docker



docker
con19
SAN FRANCISCO

Who's This Session For?

- You know some Node
- You know some Docker
- You want more Node+Docker awesomesauce

What We Gonna' Learn Bret?

- Node Dockerfile Best Practices
- Make a real-world multi-stage Dockerfile
- Build with auditing and sec scans
- Proper Node shutdown
- Node HTTP connection management

Node Dockerfiles



docker
con19
BUILD YOUR CLOUD

Every Node Sample Dockerfile

```
1 FROM node:12
2
3 EXPOSE 3000
4
5 WORKDIR /app
6
7 COPY package.json package-lock*.json ./
8
9 RUN npm install && npm cache clean --force
10
11 COPY . .
12
13 CMD ["npm", "start"]
```



Node Base Image Guidelines

- Stick to even numbered major releases
- Don't use **:latest** tag
- Start with Debian if migrating
- Use stretch (default) not jessie
- Try slim first
- Move to Alpine later, maybe



Resources

Node LTS Releases <https://github.com/nodejs/Release#release-schedule>

why does slim say not recommended: <https://github.com/nodejs/docker-node/issues/26>

default images often based on buildpacks (big): https://hub.docker.com/_/buildpack-deps

ONBUILD not recommended <https://github.com/docker-library/official-images/issues/2076>

When to use Alpine Images

- Alpine is "small" and "sec focused"
- But Debian/Ubuntu are smaller now too
- ~100MB space savings isn't significant
- Alpine has its own issues
- Alpine CVE scanning fails
- Enterprises may require CentOS or Ubuntu/Debian



<https://kubedex.com/follow-up-container-scanning-comparison/>

base image security comparison

https://docs.google.com/spreadsheets/d/1GxtCRMyKKvWUvslsGTpqUZ1wPnh1KMGha-Umkdo_xQE/edit#gid=0

Image Sizes for node/slim/alpine

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
node	12.0-slim	8651cebb80e1	23 hours ago	150MB
node	12.0	d97e1f326ca9	23 hours ago	906MB
node	12.0-alpine	80a733d0cd8c	23 hours ago	77.3MB
node	10-slim	914bfdbef6aa	4 weeks ago	143MB
node	10-stretch	64c810caf95a	4 weeks ago	899MB
node	10-jessie	5c6c62fac703	4 weeks ago	680MB

Image Sizes for node/slim/alpine

TAG	IMAGE ID	CREATED	SIZE
12.0-slim	8651cebb80e1	23 hours ago	150MB
12.0	d97e1f326ca9	23 hours ago	906MB
12.0-alpine	80a733d0cd8c	23 hours ago	77.3MB
10-slim	914bfdbef6aa	4 weeks ago	143MB
10-stretch	64c810caf95a	4 weeks ago	899MB
10-jessie	5c6c62fac703	4 weeks ago	680MB

node_modules in Images

- **Problem:** we shouldn't build images with **node_modules** from host
- Example: **node-gyp**
- **Solution:** add **node_modules** to **.dockerignore**
- copy **.gitignore**?



we have two problems when building local images with our repos that have node_modules as a sub directory

1. by default that node_modules will be copied into image
1. fix: add to .dockerignore

Least Privilege: Using **node** User

- Official node images have a node user
- But it's not used until you **USER node**
- Do this after **apt/apk** and **npm i -g**
- Do this before **npm i**
- May cause permissions issues with write access
- May require **chown node:node**



<https://docs.docker.com/engine/reference/builder/#user>

<https://docs.docker.com/engine/reference/builder/#copy>

0.Dockerfile

```
1 FROM node:12
2
3 EXPOSE 3000
4
5 WORKDIR /app
6
7 COPY package.json package-lock*.json ./
8
9 RUN npm install && npm cache clean --force
10
11 COPY . .
12
13 CMD ["npm", "start"]
```



1.Dockerfile

```
1 FROM node:12-slim
2
3 EXPOSE 3000
4
5 RUN mkdir /app && chown -R node:node /app
6
7 WORKDIR /app
8
9 USER node
10
11 COPY --chown=node:node package.json package-lock*.json ./
12
13 RUN npm install && npm cache clean --force
14
15 COPY --chown=node:node . .
16
17 CMD ["npm", "start"]
```



Process Management and Shutdown



docker
con19
2019

Node Process Management In Containers

- No need for nodemon, forever, or pm2 on server
 - We'll use nodemon in dev for file watch later
- Docker manages app start, stop, restart, healthcheck
- Node multi-thread: Docker manages multiple “replicas”
- One npm/node problem: They don't listen for proper shutdown signal by default



The Truth About The PID 1 Problem

- PID 1 (Process Identifier) is the first process in a system (or container) (AKA init)
- Init process in a container has two jobs:
 - reap zombie processes
 - pass signals to sub-processes
- Zombie not a big Node issue
- Focus on proper Node shutdown

Proper CMD for Healthy Shutdown

- Docker uses Linux signals to stop app (SIGINIT/SIGTERM/SIGKILL)
- SIGINIT/SIGTERM allow graceful stop
- npm doesn't respond to SIGINIT/SIGTERM
- node doesn't respond by default, but can with code
- Docker provides a init PID 1 replacement option

Proper Node Shutdown Options

- **Temp:** Use `--init` to fix ctrl-c for now
- **Workaround:** add `tini` to your image
- **Production:** your app captures `SIGINT` for proper exit

Example init command

- Run any node app with `--init` to handle signals (temp solution)

> `docker run --init -d nodeapp`

Example tini Dockerfile

- Add tini to your Dockerfile, then use it in CMD (permanent workaround)

> RUN apk add --no-cache tini

> ENTRYPOINT ["/sbin/tini", "--"]

> CMD ["node", "./bin/www"]

1.Dockerfile

```
1 FROM node:12-slim
2
3 EXPOSE 3000
4
5 RUN mkdir /app && chown -R node:node /app
6
7 WORKDIR /app
8
9 USER node
10
11 COPY --chown=node:node package.json package-lock*.json ./
12
13 RUN npm install && npm cache clean --force
14
15 COPY --chown=node:node . .
16
17 CMD ["npm", "start"]
```



2.Dockerfile

```
1 FROM node:12-slim
2
3 ENV TINI_VERSION v0.18.0
4 ADD https://github.com/krallin/tini/releases/download/${TINI_VERSION}/tini /tini
5 RUN chmod +x /tini
6 ENTRYPOINT ["/tini", "--"]
7
8 EXPOSE 3000
9
10 RUN mkdir /app && chown -R node:node /app
11
12 WORKDIR /app
13
14 USER node
15
16 COPY --chown=node:node package.json package-lock*.json ./
17
18 RUN npm install && npm cache clean --force
19
20 COPY --chown=node:node . .
21
22 CMD ["node", "./bin/www"]
```

Example SIGINT Capture

```
1 // quit on ctrl-c when running docker in terminal
2 process.on('SIGINT', function onSigint () {
3   console.info('Got SIGINT (aka ctrl-c in docker). Graceful shutdown ', new Date().toISOString());
4   shutdown();
5 });
6
7 // quit properly on docker stop
8 process.on('SIGTERM', function onSigterm () {
9   console.info('Got SIGTERM (docker container stop). Graceful shutdown ', new Date().toISOString());
10  shutdown();
11 });
12
13 // shut down server
14 function shutdown() {
15   server.close(function onServerClosed (err) {
16     if (err) {
17       console.error(err);
18       process.exitCode = 1;
19     }
20     process.exit();
21   });
22 }
```



```
1 // quit on ctrl-c when running docker in terminal
2 process.on('SIGINT', function onSigint () {
3     console.info('Got SIGINT (aka ctrl-c in docker). Graceful shutdown ', new Date().toIS
4     shutdown();
5 });
6
7 // quit properly on docker stop
8 process.on('SIGTERM', function onSigterm () {
9     console.info('Got SIGTERM (docker container stop). Graceful shutdown ', new Date().toIS
10    shutdown();
11 });
12
13 // shut down server
14 function shutdown() {
15     server.close(function onServerClosed (err) {
16         if (err) {
17             console.error(err);
18             process.exitCode = 1;
19         }
20         process.exit();
21     });
22 }
```



Better: Connection Tracking

- Used to track HTTP connections and send them FIN packets when Node shuts down
- > <https://github.com/hunterloftis/stoppable>

Multi-stage Builds

- Build multiple images from one file
- Those images can FROM each other
- COPY files between them
- Space + security benefits
- Great for "artifact only"
- Great for dev + test + prod

Avoiding devDependencies In Prod

- Multi-stage can solve this
- prod stages: `npm i --only=production`
- Dev stage: `npm i --only=development`
- Optional: Use `npm ci` to speed up builds
- Ensure `NODE_ENV` is set

2.Dockerfile

```
1 FROM node:12-slim
2
3 ENV TINI_VERSION v0.18.0
4 ADD https://github.com/krallin/tini/releases/download/${TINI_VERSION}/tini /tini
5 RUN chmod +x /tini
6 ENTRYPOINT ["/tini", "--"]
7
8 EXPOSE 3000
9
10 RUN mkdir /app && chown -R node:node /app
11
12 WORKDIR /app
13
14 USER node
15
16 COPY --chown=node:node package.json package-lock*.json ./
17
18 RUN npm install && npm cache clean --force
19
20 COPY --chown=node:node . .
21
22 CMD ["node", "./bin/www"]
```

3.Dockerfile

```
1 FROM node:12-slim as prod
2 ENV NODE_ENV=production
3 ENV TINI_VERSION v0.18.0
4 ADD https://github.com/krallin/tini/releases/download/${TINI_VERSION}/tini /tini
5 RUN chmod +x /tini
6 ENTRYPOINT ["/tini", "--"]
7 EXPOSE 3000
8 RUN mkdir /app && chown -R node:node /app
9 WORKDIR /app
10 USER node
11 COPY --chown=node:node package.json package-lock*.json ./
12 RUN npm ci && npm cache clean --force
13 COPY --chown=node:node . .
14 CMD ["node", "./bin/www"]
15
16 FROM prod as dev
17 ENV NODE_ENV=development
18 ENV PATH=/app/node_modules/.bin:$PATH
19 RUN npm install --only=development
20 # NOTE CHANGE ENTRYPOINT?
21 CMD ["nodemon", "./bin/www", "--inspect=0.0.0.0:9229"]
```

Building A Specific Stage

- To build dev image from dev (last) stage
> **docker build -t myapp .**
- To build prod image from prod stage
> **docker build -t myapp:prod --target prod .**

More Multi-stage: test

- Add a test stage that runs `npm test`
- Have CI build `--target test` stage before building prod
- Don't **COPY** code into dev stage
- Keep it DRY (for **COPY** and **RUN**)

3.Dockerfile

```
1 FROM node:12-slim as prod
2 ENV NODE_ENV=production
3 ENV TINI_VERSION v0.18.0
4 ADD https://github.com/krallin/tini/releases/download/${TINI_VERSION}/tini /tini
5 RUN chmod +x /tini
6 ENTRYPOINT ["/tini", "--"]
7 EXPOSE 3000
8 RUN mkdir /app && chown -R node:node /app
9 WORKDIR /app
10 USER node
11 COPY --chown=node:node package.json package-lock*.json ./
12 RUN npm ci && npm cache clean --force
13 COPY --chown=node:node . .
14 CMD ["node", "./bin/www"]
15
16 FROM prod as dev
17 ENV NODE_ENV=development
18 ENV PATH=/app/node_modules/.bin:$PATH
19 RUN npm install --only=development
20 # NOTE CHANGE ENTRYPOINT?
21 CMD ["nodemon", "./bin/www", "--inspect=0.0.0.0:9229"]
```


4.Dockerfile

```
1 FROM node:12-slim as base
2 ENV NODE_ENV=production
3 ENV TINI_VERSION v0.18.0
4 ADD https://github.com/krallin/tini/releases/download/${TINI_VERSION}/tini /tini
5 RUN chmod +x /tini
6 EXPOSE 3000
7 RUN mkdir /app && chown -R node:node /app
8 WORKDIR /app
9 USER node
10 COPY --chown=node:node package.json package-lock*.json ./
11 RUN npm ci && npm cache clean --force
12
13 FROM base as dev
14 ENV NODE_ENV=development
15 ENV PATH=/app/node_modules/.bin:$PATH
16 RUN npm install --only=development
17 CMD ["nodemon", "./bin/www", "--inspect=0.0.0.0:9229"]
18
19 FROM base as source
20 COPY --chown=node:node . .
21
22 FROM source as test
23 ENV NODE_ENV=development
24 ENV PATH=/app/node_modules/.bin:$PATH
25 COPY --from=dev /app/node_modules /app/node_modules
26 RUN eslint .
27 RUN npm test
28 CMD ["npm", "run", "test"]
29
30 FROM source as prod
31 ENTRYPOINT ["/tini", "--"]
32 CMD ["node", "./bin/www"]
```

4.Dockerfile

```
13 FROM base as dev
14 ENV NODE_ENV=development
15 ENV PATH=/app/node_modules/.bin:$PATH
16 RUN npm install --only=development
17 CMD ["nodemon", "./bin/www", "--inspect=0.0.0.0:9229"]
18
19 FROM base as source
20 COPY --chown=node:node . .
21
22 FROM source as test
23 ENV NODE_ENV=development
24 ENV PATH=/app/node_modules/.bin:$PATH
25 COPY --from=dev /app/node_modules /app/node_modules
26 RUN eslint .
27 RUN npm test
28 CMD ["npm", "run", "test"]
29
30 FROM source as prod
31 ENTRYPOINT ["/tini", "--"]
32 CMD ["node", "./bin/www"]
```

Security Scanning and Audit

- Create audit stage for optional build
- Consider **RUN npm audit**
- Consider CVE scanner
- Only report at first, no failing (most images have at least one CVE vuln)



[./multi-stage-scanning/](#)

<https://snyk.io/blog/ten-npm-security-best-practices/>

TODO recommend a few scanners

<https://kubedex.com/container-scanning/>

<https://kubedex.com/follow-up-container-scanning-comparison/>

<https://sysdig.com/blog/container-security-docker-image-scanning/>

<https://sysdig.com/blog/20-docker-security-tools/>

TODO check gareth's slides for other tools

4.Dockerfile

```
13 FROM base as dev
14 ENV NODE_ENV=development
15 ENV PATH=/app/node_modules/.bin:$PATH
16 RUN npm install --only=development
17 CMD ["nodemon", "./bin/www", "--inspect=0.0.0.0:9229"]
18
19 FROM base as source
20 COPY --chown=node:node . .
21
22 FROM source as test
23 ENV NODE_ENV=development
24 ENV PATH=/app/node_modules/.bin:$PATH
25 COPY --from=dev /app/node_modules /app/node_modules
26 RUN eslint .
27 RUN npm test
28 CMD ["npm", "run", "test"]
29
30 FROM source as prod
31 ENTRYPOINT ["/tini", "--"]
32 CMD ["node", "./bin/www"]
```

4.Dockerfile

```
19 FROM base as source
20 COPY --chown=node:node . .
21
22 FROM source as test
23 ENV NODE_ENV=development
24 ENV PATH=/app/node_modules/.bin:$PATH
25 COPY --from=dev /app/node_modules /app/node_modules
26 RUN eslint .
27 RUN npm test
28 CMD ["npm", "run", "test"]
29
30 FROM source as prod
31 ENTRYPOINT ["/tini", "--"]
32 CMD ["node", "./bin/www"]
```

5.Dockerfile

```
19 FROM base as source
20 COPY --chown=node:node . .
21
22 FROM source as test
23 ENV NODE_ENV=development
24 ENV PATH=/app/node_modules/.bin:$PATH
25 COPY --from=dev /app/node_modules /app/node_modules
26 RUN eslint .
27 RUN npm test
28 CMD ["npm", "run", "test"]
29
30 FROM test as audit
31 USER root
32 RUN npm audit --audit-level critical
33 ARG MICROSCANNER_TOKEN
34 ADD https://get.aquasec.com/microscanner /
35 RUN chmod +x /microscanner
36 RUN /microscanner $MICROSCANNER_TOKEN --continue-on-failure
37
38 FROM source as prod
39 ENTRYPOINT ["/tini", "--"]
40 CMD ["node", "./bin/www"]
```

Got Compose?



docker
con19
DOCKERS CONGRESS

Compose YAML v2 vs v3

- Myth busting: **v3 does not replace v2**
- **v2 focus:** single-node dev/test
- **v3 focus:** multi-node orchestration
- **If not using Swarm/Kubernetes, stick to v2**



Resources

<https://docs.docker.com/compose/compose-file/compose-versioning/>

<https://github.com/docker/docker.github.io/pull/7593>

Every Node Sample Compose

```
1  version: '2.4'
2
3  services:
4    node:
5      build: .
6
7    db:
8      image: postgres
```

node_modules in Bind-Mounts

- **Problem:** we can't just bind-mount **node_modules** content from host on macOS/Windows (different arch)
- **Two Potential Solutions**



1. when we bind-mount, or host node_modules will overwrite what the image built
 1. fix: move node_modules in image and prevent node_modules in host from showing up using empty volume

node_modules in Bind-Mounts

- **Solution 1, common but less flexible:**
 - Bind-mount /app which includes modules
 - You can't **docker-compose up** until you've used **docker-compose run**
 - **node_modules** on host is now only usable from container
 - Never **npm install** from host



1. when we bind-mount, or host node_modules will overwrite what the image built
 1. fix: move node_modules in image and prevent node_modules in host from showing up using empty volume

node_modules in Bind-Mounts

- **Solution 2, more complex but flexible:**
 - Move **node_modules** up a directory in Dockerfile
 - Use empty volume to hide **node_modules** on bind-mount
 - **node_modules** on host doesn't conflict

Bind-Mounting: Performance

- On Linux, bind-mounts are native
- On macOS add **delegated** write mode
- Slower in Windows, mounting across Samba/SMB
 - Consider file sync if it gets real bad
 - Or WSL + Docker



macOS

<https://blog.docker.com/2017/05/user-guided-caching-in-docker-for-mac/>

<https://docs.docker.com/docker-for-mac/osxfs/>

<https://docs.docker.com/docker-for-mac/osxfs-caching/>

Windows

<https://github.com/cweagans/docker-bg-sync>

https://www.reddit.com/r/docker/comments/8hp6v7/setting_up_docker_for_windows_and_wsl_to_work/

```
1  version: '2.4'
2
3  services:
4    node:
5      build: .
6
7    db:
8      image: postgres
```

```
1  version: '2.4'
2
3  services:
4    node:
5      build:
6        dockerfile: 5.Dockerfile
7        context: .
8        target: dev
9      volumes:
10     - ./app:delegated
11     ports:
12     - "3000:3000"
13
14   db:
15     image: postgres
```

1.docker-compose.yml



File Monitoring and Node Auto Restarts

- Use **nodemon** for compose file monitoring
- webpack-dev-server, etc. work the same
- If Windows, enable polling
- Create a **nodemon.json** for advanced workflows (bower, webpack, parcel)



<https://github.com/remy/nodemon#config-files>

Startup Order and Dependencies

- **Problem:** Multi-service apps start out of order, node might exit or cycle
- Multi-container dependencies need:
 - Name resolution (DNS)
 - Connection failure handling

Dependency Awareness

- **depends_on:** service A needs service B
- Fixes name resolution issues with
"can't resolve <service_name>"
- Only for compose, not Orch
- compose YAML v2: works with
healthchecks like a "wait for script"

```
1  version: '2.4'
2
3  services:
4    node:
5      build:
6        dockerfile: 5.Dockerfile
7        context: .
8        target: dev
9      volumes:
10     - .:/app:delegated
11     ports:
12     - "3000:3000"
13
14   db:
15     image: postgres
```

1.docker-compose.yml



2.docker-compose.yml

```
1 version: '2.4'
2
3 services:
4   node:
5     build:
6       dockerfile: 5.Dockerfile
7       context: .
8       target: dev
9     volumes:
10      - ./app:delegated
11     ports:
12      - "3000:3000"
13     depends_on:
14       db:
15         condition: service_healthy
16
17   db:
18     image: postgres
19     healthcheck:
20       test: pg_isready -U postgres -h 127.0.0.1
21       interval: 5s
```



Production Checklist

- `CMD` node directly
- Build with `.dockerignore`
- capture `SIGTERM`, properly shutdown
- `npm ci` or `npm i --only=production`
- Scan/audit/test during builds
- Healthchecks (readiness/liveness)



resources: <https://www.bretfisher.com/docker>

new course coupon: <https://www.bretfisher.com/node>

<https://www.twitter.com/BretFisher>